This morning. I woke up and meditated.

Claude read 47 emails.

Identified the five most urgent.

Drafted replies in my Gmail drafts.

Added call reminders to my calendar.

I didn't type a single word.

That's the shift. Let me show you how it works.

LIGHTNING LESSON · APRIL 23, 2026

Stop Asking Claude.
Let Claude Work.

Shankha Dey

Senior Director PM, Salesforce Data, Search, AI
Fmr. Microsoft Bing AI Platform, Cortana/Copilot, AI Search

Jason P. Yoong

Co-Founder & COO | former Amazon, VP at Dentsu, Startup ($8M seed) | Advisory Board Member

WHO'S IN THE ROOM

You told us what eats your week.

Product Design	Healthcare & Life Sciences	Consulting / Food & Bev

One person sent us this:

Morning brief · proposals · invoicing · expense tracking · filtering leads · prospect research · candidate screening

Seven tasks. Every week. By hand. Today we fix that.

300 people signed up for this session. That number tells us something.

THE PROBLEM

Claude hasn't fixed your workload. Because you're still in the loop.

You ask. It answers.

Then you copy the answer, open the app, do the work yourself. Every time.

You prompt. It drafts.

Then you paste it, format it, save it, find it later. Every time.

You close the tab.

The output disappears. Nothing runs unless you're there asking.

The problem isn't your prompts. It's that Chat only works when you're asking.

THE SHIFT

It's not about better prompts.

Chat	Cowork
You initiate every task.	You delegate once.
You wait for the response.	It runs on a schedule.
You do something with the output.	Output appears as real files on your machine.
Nothing happens when you're not there.	Runs at 6am while you're making coffee.
Output lives in a chat window.	Nothing disappears when you close a tab.

Same Claude. Different relationship.

LIVE DEMO

I typed this. Once. Three days ago.

Every morning at 6am: read my emails, identify the five most important and urgent, draft replies for each in my Gmail drafts, then check my calendar and add a reminder for any that need a call.

Here's what was waiting for me this morning:

📥 **Gmail Drafts**

5 replies drafted

📅 **Calendar**

3 call reminders added

📄 **Morning brief**

saved to /outputs

THE FRAMEWORK

How to delegate to Cowork. Four fields.

01 END STATE

What done looks like. Not what to do — what the finished result is.

Most people describe the process. Describe the result instead.

02 CONTEXT

What Claude can't see or infer from the files.

e.g. 'These are rough notes, not final decisions.'

03 CONSTRAINTS

What it must never do.

Always include: do not delete anything unless I explicitly ask.

04 OUTPUT

What file to create and where to save it.

e.g. 'Save as q1-update.docx in /outputs.'

How do you know it worked?

The morning email job — offline evals

DETERMINISTIC CHECKS

- ✓ Exactly 5 drafts created in Gmail?
- ✓ All 5 are in Drafts, not Sent?
- ✓ Calendar reminders added for emails that needed calls?
- ✓ Task completed before 6:05am?
- ✓ Morning brief file saved to /outputs?

LLM-AS-JUDGE CHECKS

Did it prioritise correctly?
Score: would a human reviewer agree these were the 5 most urgent?

Do drafts match my voice?
Score: does each reply sound like I wrote it?

Are replies relevant?
Score: does the reply actually address what the email asked?

Define what good looks like before the job runs. Then the eval scores itself.

How do you improve it?

The morning email job — A/B tests

TEST 1 · PRIORITISATION CRITERIA

A: "identify the five most important and urgent" → **B: "identify the five that require action before 5pm today"**

Metric: LLM judge scores — did adding a time constraint improve prioritisation accuracy?

TEST 2 · VOICE CONTEXT

A: run without brand-voice.md loaded → **B: run with brand-voice.md + 3 writing samples**

Metric: human review — do drafts in Condition B sound more like me?

TEST 3 · OUTPUT SPECIFICATION

A: "add a reminder for any that need a call" → **B: "add a 15-minute calendar block titled [Sender] — follow-up call, with the email subject in the notes"**

Metric: deterministic — are B reminders more usable without editing?

Run A and B as two different jobs and assess the results by each. The diff is your improvement signal

WHAT'S NEXT

This was the 30-minute version.

May 9 is the 5 hour workshop.

You build three automations for your actual job — not demos.

You write your own Skill for your most repetitive task.

You get live instructor feedback on every piece you build.

You leave with a certificate and a working system.

50 seats · Founding cohort pricing · Link in chat

maven.com/level-up-leadership-team/claude-cowork-certification

April 30 — one more free session: Your First Automated Morning (Built Live)